



Digital Systems for the MITRA

(GPU Computing)

Submitted by

Ruben Anderson Louis

Project submitted in the partial fulfillment for the degree of

BSc (Hons) Physics with Computing

UNIVERSITY OF MAURITIUS

FACULTY OF SCIENCE

DEPARTMENT OF PHYSICS

March, 2015

To my parents

Declaration of Authorship

I, AUTHOR NAME, declare that this thesis titled, 'THESIS TITLE' and the work presented in it are my own. I confirm that:

- This work was done wholly or mainly while in candidature for a research degree at this University.
- Where any part of this thesis has previously been submitted for a degree or any other qualification at this University or any other institution, this has been clearly stated.
- Where I have consulted the published work of others, this is always clearly attributed.
- Where I have quoted from the work of others, the source is always given. With the exception of such quotations, this thesis is entirely my own work.
- I have acknowledged all main sources of help.
- Where the thesis is based on work done by myself jointly with others, I have made clear exactly what was done by others and what I have contributed myself.

Signed:

Date:

“An opening quote will be put here.”

The author’s name will be there.

UNIVERSITY OF MAURITIUS

Abstract

Faculty of Science
Department of Physics

BSc (Hons) Physics with computing

Digital Systems for the MITRA (GPU Computing)

by [Ruben Anderson Louis](#)

... This dissertation describes the development, and testing of software and general purpose computing hardware as part of the integration of Digital Systems for the Multi-frequency Interferometry for Astronomy (MITRA) array. High performance software correlation was implemented with the DiFX (Distributed FX) software by Adam Deller, making it an off the shelf, user-friendly and custom documented software further possibilities offered by General-purpose computing on Graphics Processing Units (GPGPU) were implemented, developed, and tested taking advantage of its advance parallelism, and performance with typical speedups of ... compared to ... with the DiFX software ...

Acknowledgements

On the long road to the accomplishment of my undergraduate project and thesis I wish to acknowledge the following, first I am very grateful to my supervisor, Dr. G. K. Beeharry, for the opportunity he gave me to work on the topic, he has been an inspirational lecturer who encouraged me to pursue in the field of scientific computing, also most importantly I wish to thank him for his guidance, useful suggestions during the course of this project and his reviews, remarks and comments during the writing process of the thesis. I express my thanks to the lecturers of the Department of physics who during those 3 years here at the University of Mauritius contributed to build up the physicist I am today.

My special thanks goes to the MPhil student, Mr. Vinand Prayag, who was the first to trust me to work with him on the specific topic of software correlation, he was always there to talk and help when I needed advices, both on my project and my thesis.

I am also thankful to the technical staff of the Mauritius Radio Telescope and that of the department of physics for their help and guidance during the practical sessions and implementation of my project.

Extra credits goes to my classmates with whom we struggled with the BSc (Hons) Physics course, specially those of the computing branch who will recognise themselves, where the last two years together were a lot of fun and made my stay at the university a lot nicer.

I would also like to thank those closest to me for their support and my family who has always been behind me and helped me greatly to continue in the field of Physics at university level.

Last but not least as I am a faithful believer, and I believe that nothing happens without a reason I wish to thank God for everything, and making all of this possible whichever way it may have happened.

Contents

Declaration of Authorship	i
Abstract	iii
Acknowledgements	iv
List of Figures	vii
List of Tables	viii
Abbreviations	ix
Physical Constants	x
Symbols	xi
1 Introduction	1
1.1 Digital Systems for the MITRA	1
1.2 Aims and objectives	1
1.3 Outline	2
2 Literature Review	3
2.1 Mr. Vinand Prayag Lit. Rev. Plan advice	3
2.2 [Cheesy] Preliminary gather up of abstract or similar	3
3 Synthesis Imaging	13
4 Distributed FX Correlation - The DiFX software	14
5 Conclusions and future works	15
5.1 Conclusions	15
5.2 Future works	15
A Background, mathematical derivations	16

B	DiFX - Custom documentation	17
B.1	DiFX - Preliminary setup	17
C	The Git repository	21
D	Progress Log	22

List of Figures

List of Tables

Abbreviations

EM	E lectrom m agnetic
FFT	F ast F ourier T ransform
GPU	G raphics P rocessing U nit
GPGPU	G eneral P urpose C omputing on G raphics P rocessing U nits
MEM	M aximum E ntropy M ethod
MITRA	M ultifrequency I nterferometry T elescope for R adio A stronomy
MRT	M auritius R adio T elescope
NNLS	n on- n egative least square a lgorithm
R.A.	R adio A stronomy
SDR	S oftware D efined R adio
VLA	V ery L arge A rray (Radiotelescope)

Physical Constants

$$\text{Speed of Light } c = 2.997\,924\,58 \times 10^8 \text{ ms}^{-1} \text{ (exact)}$$

Symbols

a	distance	m
P	power	W (Js ⁻¹)
ω	angular frequency	rads ⁻¹

Chapter 1

Introduction

1.1 Digital Systems for the MITRA

...

1.2 Aims and objectives

This project is aimed at the development, and testing of software and general purpose computing hardware for the Multi-frequency Interferometry for Astronomy (MITRA) array project. In so doing it will pave the way for a higher integration of digital systems as opposed to specialised hardware in the synthesis imaging process. The particular application which was focused on was that of high performance software correlation by the use of parallel computing. The central objective was to implement the DiFX (Distributed FX) software by Adam Deller for the MITRA by making it an off the shelf, user-friendly and custom documented software. The second part of the project was to explore the further possibilities opened up by General-purpose computing on Graphics Processing Units (GPGPU) to implement, develop, and test code taking advantage of its advance parallelism, and performance. I hope that the that the work on this particular project will not stay still and that this dissertation will inspire other people to continue to the work, thus from that initiative follows naturally the creation a public git repository: https://github.com/Benzy-Louis/MITRA_FX-CPU-GPU.git where those interested to contribute to the project can continue the development, work on the source code, and improve software and its documentation on their own or contribute for the same to the repository, those interested are freely encouraged to mail me at louis.ruben@gmail.com if they want to contribute or have questions, suggestions on the project or on this dissertation and an appendix is written for the purpose of the use of the git repository ...

1.3 Outline

The dissertation is composed of . . . main chapters where . . . It is structured as follows *chapter 1* is an introduction to . . . Then *chapter 2* . . .

Chapter 2

Literature Review

[Everything (codes,text, formatting) in this chapter is cheesy for the sake of having a quick overview while preparing the real literature review]

2.1 Mr. Vinand Prayag Lit. Rev. Plan advice

The past

Hardware Correlator -> e.g. MRT room size, hardware . . .

Evolution

Better processing power CPU made it feasible to replace hardware based system by software based ones. A rapid evolution of software defined radio (SDR). Direct applicability to radio astronomy.

Examples

Now we have what we call software telescopes such as the LOFAR powered by IBM supercomputers.

Recently SWINBURNE built a general purpose software correlator the DiFX,

Pros. and cons. of DiFX.

2.2 [Cheesy] Preliminary gather up of abstract or similar

2008

Jheengut - Software Correlation

Software correlation is seen to replace digital correlation as a step forward in removing the excessive cost for dedicated hardware in the near future. A software correlator for S. Ord, L. Greenhill, R. Wayth, D. Mitchell, K. Dale, H. Pfister, R. G. Edgar - GPUs for data processing in the MWA

The MWA is a next-generation radio interferometer under construction in remote Western Australia. The data rate from the correlator makes storing the raw data infeasible, so the data must be processed in real-time. The processing task is of order 10 TFLOPs^{-1} . The remote location of the MWA limits the power that can be allocated to computing. We describe the design and implementation of elements of the MWA real-time data processing system which leverage the computing abilities of modern graphics processing units (GPUs). The matrix algebra and texture mapping capabilities of GPUs are well suited to the majority of tasks involved in real-time calibration and imaging. Considerable performance advantages over a conventional CPU-based reference implementation are obtained.

Chris Harris Karen Haines Lister Staveley-Smith -
GPU Accelerated Radio Astronomy Signal Convolution

The increasing array size of radio astronomy interferometers is causing the associated computation to scale quadratically with the number of array signals. Consequently, efficient usage of alternate processing architectures should be explored in order to meet this computational challenge. Affordable parallel processors have been made available to the general scientific community in the form of the commodity graphics card. This work investigates the use of the Graphics Processing Unit (GPU) in the parallelisation of the combined conjugate multiply and accumulation stage of a correlator for a radio astronomy array. Using NVIDIA's Compute Unified Device Architecture, our testing shows processing speeds from one to two orders of magnitude faster than a Central Processing Unit (CPU) approach.

Andrew Woods, Michael Inggs and Alan Langman -
Accelerating a Software Radio Astronomy Correlator using FPGA co-processors

This article presents and characterises our work on accelerating a software radio astronomy correlator using reconfigurable computing (RC) hardware. Radio astronomy correlation is an embarrassingly parallel signal processing application, which is used heavily in radio astronomy for imaging and other astronomical measurements. Radio astronomy correlators typically operate on huge data sets and often require real-time processing, as storage of raw data is impractical - resulting in substantial computational requirement. Currently FPGAs are the preferred processing architecture used in modern large radio astronomy correlators [1] and perform well on the types of DSP functions that correlators perform. In this paper we set out to accelerate the DiFX (Distributed FX) correlator, a software correlator, using FPGA reconfigurable computing hardware. — hoping to inherit some of the advantages that larger production FPGA correlators have over software.

2009

NVIDIA -

NVIDIA's Next Generation CUDA™ Compute Architecture: Fermi™

Adam Trevis Deller -

Precision VLBI astrometry: Instrumentation, algorithms and pulsar parallax determination

This thesis describes the development of DiFX, the first general-purpose software correlator for radio interferometry, and its use with the Australian Long Baseline Array to complete the largest Very Long Baseline Interferometry (VLBI) pulsar astrometry program undertaken to date in the Southern Hemisphere. This two year astrometry program has resulted in the measurement of seven new pulsar parallaxes, which has more than trebled the number of measured VLBI pulsar parallaxes in the Southern Hemisphere. These measurements included a determination of the distance and transverse velocity of PSR J0437–4715 with better than 1% accuracy, enabling improved tests of General Relativity; the first significant measurement of parallax for the famous double pulsar system PSR J0737–3039A/B, which will allow tests of General Relativity in this system to proceed to the 0.01% level and also offers insights into its formation and high-energy emission; and a factor of four revision to the estimated distance of PSR J0630–2834, which had previously appeared to possess extremely unusual x-ray emission characteristics. Additionally, the ensemble of refined distance and transverse velocity estimates have enabled a widely applicable improvement in knowledge of pulsar luminosities in several wavebands and the Galactic electron distribution at southern latitudes. Finally, the DiFX software correlator developed to enable this science has been extensively tested and verified against three existing hardware correlators, and is now an integral part of the upgraded Long Baseline Array Major National Research Facility used by astronomers throughout Australia and the world; furthermore, it has been selected to facilitate a major upgrade of the world's only full-time VLBI instrument, the Very Long Baseline Array operated by the National Radio Astronomy Observatory in the US.

Rob V. van Nieuwpoort, John W. Romein -

Using Many-Core Hardware to Correlate Radio Astronomy Signals

A recent development in radio astronomy is to replace traditional dishes with many small antennas. The signals are combined to form one large, virtual telescope. The enormous data streams are crosscorrelated to filter out noise. This is especially challenging, since the computational demands grow quadratically with the number of data streams. Moreover, the correlator is not only computationally intensive, but also very I/O intensive. The LOFAR telescope, for instance, will produce over 100 terabytes per day. The future SKA telescope will even require in the order of exaflops, and petabits/s of I/O. A recent trend is to correlate in software instead of dedicated hardware. This is done to increase flexibility and to reduce development efforts. Examples include e-VLBI and LOFAR. In this paper, we evaluate the correlator algorithm on multi-core CPUs and many-core

architectures, such as NVIDIA and ATI GPUs, and the Cell/B.E. The correlator is a streaming, real-time application, and is much more I/O intensive than applications that are typically implemented on many-core hardware today. We compare with the LOFAR production correlator on an IBM Blue Gene/P supercomputer. We investigate performance, power efficiency, and programmability. We identify several important architectural problems which cause architectures to perform suboptimally. Our findings are applicable to data-intensive applications in general. The results show that the processing power and memory bandwidth of current GPUs are highly imbalanced for correlation purposes. While the production correlator on the Blue Gene/P achieves a superb 96% of the theoretical peak performance, this is only 14% on ATI GPUs, and 26% on NVIDIA GPUs. The Cell/B.E. processor, in contrast, achieves an excellent 92%. We found that the Cell/B.E. is also the most energy-efficient solution, it runs the correlator 5-7 times more energy efficiently than the Blue Gene/P. The research presented is an important pathfinder for next-generation telescopes.

2010

Andrew Woods, Michael Inggs and Alan Langman -

Accelerating a Software Radio Astronomy Correlator using FPGA co-processors

This thesis attempts to accelerate compute intensive sections of a frequency domain radio astronomy correlator using dedicated co-processors. Two co-processor implementations were made independently with one using reconfigurable hardware (Xilinx Virtex 4LX100) and the other uses a graphics processor (Nvidia 9800GT). The objective of a radio astronomy correlator is to compute the complex valued correlation products for each baseline which can be used to reconstruct the sky's radio brightness distribution. Radio astronomy correlators have huge computation demands and this dissertation focuses on the computational aspects of correlation, concentrating on the X-engine stage of the correlator. Although correlation is an extremely compute intensive process, it does not necessarily require custom hardware. This is especially true for older correlators or VLBI experiments, where the processing and I/O requirements can be satisfied by commodity processors in software. Discrete software co-processors like GPUs and FPGAs are an attractive option to accelerate software correlation, potentially offering better FLOPS/watt and FLOPS/\$ performance. In this dissertation we describe the acceleration of the X-engine stage of a correlator on a CUDA GPU and an FPGA. We compare the co-processors' performance with a CPU software correlator implementation in a range of different benchmarks. Speedups of 7x and 12.5x were achieved on the FPGA and GPU correlator implementations respectively. Although both implementations achieved speedups and better power utilisation than the CPU implementation, the GPU implementation produced better performance in a shorter development time than the FPGA. The FPGA implementation was hampered by the development tools and the slow PCI-X bus, which is used to communicate with the host. Additionally, the Virtex 4 LX100 FPGA was released two years before the Nvidia G80 GPU and so is more behind the current technologies. However, the FPGA does have an advantage in terms of power efficiency, but power consumption is only a concern for large compute clusters. We found that using GPUs was the better option to

accelerate small-scale software X-engine correlation than the Virtex 4 FPGA.

Nicolas PLATEL - Implémentation d'un corrélateur sur une carte GPU.

Le but de ce projet est d'implémenter un corrélateur de type FX en Software sur une carte GPU de la marque NVIDIA. Ce corrélateur permettra aux étudiants de finir la construction du télescope étudié et obtenir des images du ciel. Il a pour but également de donner quelques notions sur Cuda aux étudiants le désirants. Dans le but d'étudier des phénomènes physiques connus et de valider mon projet, des résultats expérimentaux ont été effectués grâce aux antennes et au récepteur numérique créé précédemment. Enfin, une interface graphique a été créée pour faciliter l'utilisation du corrélateur à l'utilisateur.

Hobiger T., Kimura M., Takefuji K, Oyama T., Koyama Y., Kondo T., Gotoh T., Amagai J.

GPU based software correlators-perspectives for VLBI2010

Caused by historical separation and driven by the requirements of the PC gaming industry, Graphics Processing Units (GPUs) have evolved to massive parallel processing systems which entered the area of non-graphic related applications. Although a single processing core on the GPU is much slower and provides less functionality than its counterpart on the CPU, the huge number of these small processing entities outperforms the classical processors when the application can be parallelized. Thus, in recent years various radio astronomical projects have started to make use of this technology either to realize the correlator on this platform or to establish the post-processing pipeline with GPUs. Therefore, the feasibility of GPUs as a choice for a VLBI correlator is being investigated, including pros and cons of this technology. Additionally, a GPU based software correlator will be reviewed with respect to energy consumption/GFlop/sec and cost/GFlop/sec.

Patrick Brandt, Ron Duplain, Paul Demorest, Randy McCullough, Scott Ransom, Jason Ray

Heterogeneous real-time computing in radio astronomy

Modern computer architectures suited for general purpose computing are often not the best choice for either I/O-bound or compute-bound problems. Sometimes the best choice is not to choose a single architecture, but to take advantage of the best characteristics of different computer architectures to solve your problems. This paper examines the tradeoffs between using computer systems based on the ubiquitous X86 Central Processing Units (CPU's), Field Programmable Gate Array (FPGA) based signal processors, and Graphical Processing Units (GPU's). We will show how a heterogeneous system can be produced that blends the best of each of these technologies into a real-time signal processing system. FPGA's tightly coupled to analog-to-digital converters connect the instrument to the telescope and supply the first level of computing to the system. These FPGA's are coupled to other FPGA's to continue to provide highly efficient processing power.

Data is then packaged up and shipped over fast networks to a cluster of general purpose computers equipped with GPU's, which are used for floating-point intensive computation. Finally, the data is handled by the CPU and written to disk, or further processed. Each of the elements in the system has been chosen for its specific characteristics and the role it can play in creating a system that does the most for the least, in terms of power, space, and money.

GINOURIE Sabera Bibi

A prototype front-end and back-end receiver system for radioastronomy

The first part of the project consisted of designing and building a front-end and back-end system for radioastronomy. Eight Log-periodic dipole antennas (available at the MRT) were used for the front-end. In the second part, a new data acquisition card was used. This card was studied and programmed before used. The card was tested several times in order to check whether the analog data were digitised. Next, the whole system was tested and observations were carried out. Celestial objects like Virgo A and Centaurus A were successfully observed.

2011

V. K. Veligatla, P. Labropoulos, L. V. E. Koopmans -

Adaptive Beam-forming for Radio Astronomy On GPU

The LOFAR radio telescope consists of tens of thousands of dipole antennas that combine their signals to operate as a single large radio telescope. The truly innovative aspect of this new telescope is that its pointing system is not mechanical. It is steered by combining the electric signals from different elements using advanced beam-forming software. Imaging software is one of the important aspects of processing the high-volume data streams produced by LOFAR, and is one of the best places to use GPUs to achieve processing speed. We were able to achieve up to 30 times performance gain compared to the CPU implementation in novel, computationally intensive techniques such as the Minimum Variance Distortionless Response (MVDR). We have gained 5-6 times speed-up compared to the CPU implementation for standard imaging algorithms.

M. A. Clark, P. C. La Plante, L. J. Greenhill -

Accelerating Radio Astronomy Cross-Correlation with Graphics Processing Units

We present a highly parallel implementation of the cross-correlation of time-series data using graphics processing units (GPUs), which is scalable to hundreds of independent inputs and suitable for the processing of signals from "Large-N" arrays of many radio antennas. The computational part of the algorithm, the X-engine, is implemented efficiently on Nvidia's Fermi architecture, sustaining up to 79% of the peak single precision floating-point throughput. We compare performance obtained for hardware- and software-managed caches, observing significantly better performance for the latter. The high performance reported involves use of a multi-level data tiling strategy in memory and use of a pipelined algorithm with simultaneous computation and transfer

of data from host to device memory. The speed of code development, flexibility, and low cost of the GPU implementations compared to ASIC and FPGA implementations have the potential to greatly shorten the cycle of correlator development and deployment, for cases where some power consumption penalty can be tolerated.

2012

NVIDIA -

NVIDIA's Next Generation CUDA™ Compute Architecture: Kepler™ GK110

V. K. Veligatla, P. Labropoulos, L. V. E. Koopmans -

Adaptive Beam-forming for Radio Astronomy On GPU

The LOFAR radio telescope consists of tens of thousands of dipole antennas that combine their signals to operate as a single large radio telescope. The truly innovative aspect of this new telescope is that its pointing system is not mechanical. It is steered by combining the electric signals from different elements using advanced beam-forming software. Imaging software is one of the important aspects of processing the high-volume data streams produced by LOFAR, and is one of the best places to use GPUs to achieve processing speed. We were able to achieve up to 30 times performance gain compared to the CPU implementation in novel, computationally intensive techniques such as the Minimum Variance Distortionless Response (MVDR). We have gained 5-6 times speed-up compared to the CPU implementation for standard imaging algorithms.

John W. Romein -

An Efficient Work-Distribution Strategy for Gridding Radio-Telescope Data on GPUs

This paper presents a novel work-distribution strategy for GPUs, that efficiently convolves radio-telescope data onto a grid, one of the most time-consuming processing steps to create a sky image. Unlike existing work-distribution strategies, this strategy keeps the number of device-memory accesses low, without incurring the overhead from sorting or searching within telescope data. Performance measurements show that the strategy is an order of magnitude faster than existing accelerator-based gridders. We compare CUDA and OpenCL performance for multiple platforms. Also, we report very good multi-GPU scaling properties on a system with eight GPUs, and show that our prototype implementation is highly energy efficient. Finally, we describe how a unique property of GPUs, fast texture interpolation, can be used as a potential way to improve image quality.

Alessio Sclocco, Ana Lucia Varbanescu, Jan David Mol, Rob V. van Nieuwpoort -

Radio Astronomy Beam Forming on Many-Core Architectures

Traditional radio telescopes use large steel dishes to observe radio sources. The largest radio

telescope in the world, LOFAR, uses tens of thousands of fixed, omnidirectional antennas instead, a novel design that promises ground-breaking research in astronomy. Where traditional telescopes use custom-built hardware, LOFAR uses software to do signal processing in real time. This leads to an instrument that is inherently more flexible. However, the enormous data rates and processing requirements (tens to hundreds of teraflops) make this extremely challenging. The next-generation telescope, the SKA, will require exaflops. Unlike traditional instruments, LOFAR and SKA can observe in hundreds of directions simultaneously, using beam forming. This is useful, for example, to search the sky for pulsars (i.e. rapidly rotating highly magnetized neutron stars). Beam forming is an important technique in signal processing: it is also used in WIFI and 4G cellular networks, radar systems, and health-care microwave imaging instruments. We propose the use of many-core architectures, such as 48- core CPU systems and Graphics Processing Units (GPUs), to accelerate beam forming. We use two different frameworks for GPUs, CUDA and OpenCL, and present results for hardware from different vendors (i.e. AMD and NVIDIA). Additionally, we implement the LOFAR beam former on multi-core CPUs, using OpenMP with SSE vector instructions. We use autotuning to support different architectures and implementation frameworks, achieving both platform and performance portability. Finally, we compare our results with the production implementation, written in assembly and running on an IBM Blue Gene/P supercomputer. We compare both computational and power efficiency, since power usage is one of the fundamental challenges modern radio telescopes face. Compared to the production implementation, our auto-tuned beam former is 45–50 times faster on GPUs, and 2–8 times more power efficient. Our experimental results lead to the conclusion that GPUs are an attractive solution to accelerate beam forming.

V.Vamsi Krishna, Dr. Panos Labropoulos, Prof. Leon V.E. Koopmans -
GPU's for Radio Imaging

- Signals from Sources (e.g. galaxies)
- Next Gen Antennas (e.g. LOFAR, SKA, ...)
- Image acquired after Processing (RFI elimination, Calibration).

Mike Clark with Lincoln Greenhill and Paul LaPlante -
Accelerating Radio Astronomy Cross-Correlation Beyond 1 Tflops Using Fermi

2013

Dominique Ingala -Durban University of Technology
An Overview of the MITRA Radio Telescope Signal Chain
Project proposal: by Prof. Girish Beeharry
Locations: UoM and DUT
Operational frequency range: 200 – 800 MHz
Antenna type: Dual Polarized LPDA

2 arrays (8 LPDA's per array)

Front-End RF and IF stage: Analog components

Digital Back-End: Software Defined Radio

Fiber optic links

Cost effective design

Harshavardhan Reddy Suda, Pradeep Kumar Gupta -

Powering Real-time Radio Astronomy Signal Processing with GPUs. Design of a GPU based real-time backend for the upgraded GMRT

Nitisha Pirthee -

Digital back end for MITRA prototype

In the first part of the project, USRP1 was used on GNU radio. A log periodic antenna was connected to one channel of the USRP and the expected peaks were observed. The second channel was not operational when tested. An array of sixteen channels was used as front end. A PCI-ADC card already available at MRT was used to do data acquisition. The program for data acquisition was improved. The card was tested several times and observations were carried out. Celestial objects like CAS A and Pictor A were successfully observed.

Harshavardhan Reddy Suda, Pradeep Kumar Gupta -

Powering Real-time Radio Astronomy Signal Processing with GPUs. Design of a GPU based real-time backend for the upgraded GMRT

Ben Barsdell, Mike Clark, Lincoln Greenhill, Jonathon Kocz -

ACCELERATING RADIO ASTRONOMY CROSS-CORRELATION USING THE KEPLER ARCHITECTURE

Kepler GK110 optimisation

2014

Ben Barsdell, Mike Clark, Lincoln Greenhill, Jonathon Kocz -

PETASCALE CROSS-CORRELATION

Amr H. Hassan, Christopher Fluke, David Barnes, Virginia Kilborn -

Astronomical "Big Data" Analysis and Visualization

Alex Bogert, John Holdener, and Nicholas Smith -

Interactive Visualization of Astrophysical Data

Yt is an analysis and visualization system for astrophysical volumetric data that is openly developed and freely available. At its core, yt provides a method of describing physical rather than computational objects inside an astrophysical simulation. yt provides methods for selecting regions, applying analysis to regions, visualizing (including volume rendering, projections, slices, phase plots) and exporting data to external analysis packages.

S. Bhatnagar, P. K. Gupta, M. Clark -

GPU based imager for radio astronomy

Mario Guillaume CECILE -

Enhancement of some computational physics algorithms using Parallel Computing and the Graphical Processing Unit

Scientific computing has become an important method for testing and improving current scientific models and theories. Recent developments in computer architecture have helped to study more complex systems using High Performance Computing (HPC). In this project, Grain Growth simulation and the soft-sphere Discrete Element Method (DEM) are enhanced to be able to consider larger matrix sites in the Grain Growth simulation and a large number of particles in the DEM model. Parallel computing using Message Passing Interface (MPI) is used as well as CUDA for programming on the GPU. For the Grain Growth simulation, effects of foreign particles are investigated while for the DEM model, free-falling of particles in a packed bed is studied. The results presented in this work help to give further understanding about the physics involved behind both the Grain Growth and DEM. This work furthermore demonstrates how the use of parallel processing can help scientists to enhance their code.

NVIDIA -

NVIDIA GeForce GTX 750 Ti Featuring First-Generation Maxwell GPU Technology, Designed for Extreme Performance per Watt

NVIDIA GeForce GTX 980 Featuring Maxwell, The Most Advanced GPU Ever Made.

Chapter 3

Synthesis Imaging

...

Chapter 4

Distributed FX Correlation - The DiFX software

...

Chapter 5

Conclusions and future works

5.1 Conclusions

...

5.2 Future works

...

Appendix A

Background, mathematical derivations

...

Appendix B

DiFX - Custom documentation

B.1 DiFX - Preliminary setup

[Vague content just to fill in, following. . .]

To install the diFX-2.3 software correlator do the following things,

1. Download the source of the software somewhere using the following command

```
svn co https://svn.atnf.csiro.au/difx/master_tags/DiFX-2.3
```

2. Install IPP, i.e.

```
>> cd $IPP_PATH
>> chmod +x install.sh
>> ./install.sh
```

3. Install PGPLOT, following the instructions of the following website:

<http://pendientedemigracion.ucm.es/info/Astrof/software/howto/howto-pgplot.html>

Do the following,

```
>> cd /usr/local/src
>> mv ~/Downloads/pgplot5.2.tar.gz .
>> tar zxvf pgplot5.2.tar.gz
>> mkdir /usr/local/pgplot
>> cd /usr/local/pgplot
>> cp /usr/local/src/pgplot/drivers.list
>> /usr/local/src/pgplot/makemake
    /usr/local/src/pgplot linux g77_gcc_aout
```

Edit the file makefile

```
>> sudo gedit makefile &
```

Change

```
FCOMPL=g77
# to
FCOMPL=gfortran
```

Then save, and compile

```
>> make
>> make cpg
>> make clean
```

Export the paths

```
>> export PGPLOT_DIR=/usr/local/pgplot
>> export PGPLOT_DEV=/Xserve
```

4. Install OpenMPI also.
5. Then go back to install diFX-2.3, edit the setup.bash file,

```
>> cd $DIFX_ROOT
>> sudo gedit setup.bash &
```

6. Change the following paths environment variables, an example here for my setup,

```
##### ROOT PATHS #####
export DIFXROOT=/home/ruben/Final_Year_Project/DiFX/DiFX-2.3
export DIFX_PREFIX=$DIFXROOT
export PGPLOTDIR=/usr/local/pgplot/
export IPPROOT=/opt/intel/ipp/bin/

##### COMPILER #####
export MPICXX=/usr/local/bin/mpicxx

# ALSO UNCOMMENT THESE THINGS ALSO

##### PORTS FOR DIFXMESSAGE #####
# Uncomment these to enable DIFX_MESSAGES
```

```
export DIFX_MESSAGE_GROUP=224.2.2.1
export DIFX_MESSAGE_PORT=50201
export DIFX_BINARY_GROUP=224.2.2.1
export DIFX_BINARY_PORT=50202
```

7. use the geniepc script with input the path /opt/intel

```
$DIFXROOT/applications/difx_monitor
```

8. Then we can start the installation of DiFX-2.3, go back to the root folder of diFX, and do the following,

```
>> source setup.bash
>> ./install-difx
```

9. Now for the preliminary tests I have not automated the scripts which would allow me to run diFX without having to send out preliminary parameters first. To be able to use the program one must set up the RPC. So the rpcbind package has to be installed.

10. To make rpcbind work,

```
>> sudo -i service portmap stop
>> sudo -i rpcbind -i -w
>> sudo -i service portmap start
```

Then we can start the calculation server and check it for the host,

```
>> startCalcServer
>> checkCalcServer 127.0.0.1
# Setting up the CALC_SERVER environment variable to the local host
>> export CALC_SERVER=127.0.0.1
```

To see the processes, open another terminal go to the root directory of diFX, source setup.bash and then run

It will display the processes taking place in the calculation, I attached a text file with that info.

```
>> errormon2
```

11. When all this is done, one can try the example files, for the RDV70 data,


```
>> vex2difx example.v2d  
>> calcif2 example_1.calc  
>> mpirun -np 8 mpifxcorr example_1.input
```

12. When the correlation is done, you can get the FITS-LDI file from it in the following way,

```
>> difx2fits example_1.difx
```

Appendix C

The Git repository

...

Appendix D

Progress Log

...